



Introduction of MariaDB

26th September 2017

GOTO Satoru
Customer Solutions Engineer

What is MariaDB ?

The Soul of Open Source


“

MariaDB was created to preserve openness and community, so that we can push ahead faster with the capabilities for tomorrow's applications.

”

Michael “Monty” Widenius

Founder & CTO of MariaDB



History of MySQL & MariaDB



- 1981 Unireg (base of MySQL code)
- 1994 Added SQL interface and renamed it MySQL
- 1995 **MySQL released** under dual licensing
- 2005 Oracle acquired InnoDB
- 2008 Sun acquired MySQL AB for 1 billion USD
- 2009 Monty & others left Sun in Feb 2009
to work on Maria engine in Monty Program Ab
- 2009 Oracle acquired Sun
- 2012 **MariaDB foundation** was created
- 2013 MariaDB started replacing MySQL in most
Linux distributions
- 2013 Monty Program Ab merged with SkySQL
- 2014 SkySQL Ab renamed to **MariaDB Corporation**

History of MariaDB



Feb 2010 MariaDB 5.1

Nov 2010 MariaDB 5.2

Apr 2012 MariaDB 5.3

Feb 2012 MariaDB 5.5

Mar 2014 MariaDB 10.0

Oct 2015 MariaDB 10.1

May 2017 MariaDB 10.2.6 GA

Aug 2017 **MariaDB 10.3.1 alpha**

<https://downloads.mariadb.org/mariadb/+releases/>
<https://github.com/MariaDB/server/releases>

MariaDB vs. MySQL

MariaDB vs. MySQL

More Storage Engines

- **ColumnStore** column oriented engine for Data warehousing(DWH)
- **Spider** in 10.0 and later
- **MyRocks** with great compression, in 10.2
- **Aria** MyISAM replacement with better caching
- **FederatedX** drop-in replacement for Federated
- **OQGRAPH** new in 5.2 - Open Query GRAPH
- **SphinxSE** new in 5.2
- **TokuDB** in 5.5 and later -
- **CONNECT** in 10.0 and later
- **SEQUENCE** in 10.0 and later
- **Cassandra** in 10.0

<https://mariadb.com/kb/en/library/mariadb-vs-mysql-features/>

MariaDB vs. MySQL

Performance Improvements

- Query Optimizer enhancements
- Faster and safer replication
- Character set conversions
- **Thread Pool**
 - allows 200,000+ connections and with a notable speed improvement when using many connections
 - w/ MySQL **only available with Enterprise version**

<https://mariadb.com/kb/en/library/mariadb-vs-mysql-features/>

MariaDB vs. MySQL

Extensions & New Features

- **Window functions** new in 10.2
- DECIMAL from 30 to 38 new in 10.2
- **Recursive CTE** new in 10.2 (Common Table Expressions)
- CHECK CONSTRAINT new in 10.2
- DEFAULT new in 10.2, for BLOB and TEXT
- Added catchall for list partitions new in 10.2
- EXECUTE IMMEDIATE new in 10.2 (Oracle-style)
- JSON functions new in 10.2
- ms Precision in Processlist
- Table Elimination
- Virtual Columns new in 5.2

MariaDB vs. MySQL

Extensions & New Features

- Extended User Statistics new in 5.2
- KILL all queries for a user new in 5.3
- KILL QUERY ID terminates the query by query_id, leaving the connection intact — new in 10.0.5
- Pluggable Authentication new in 5.2
- Enhancements to INFORMATION SCHEMA.PLUGINS table — new in 5.2
- Group commit for the binary log new in 5.3
- Added --rewrite-db mysqlbinlog option to change the used database — new in 5.2
- Progress reporting for ALTER TABLE and LOAD DATA INFILE — new in 5.3
- Faster joins and subqueries new in 5.3
- HandlerSocket and faster HANDLER calls new in 5.3
- Dynamic Columns support new in 5.3
- GIS Functionality new in 5.3

MariaDB vs. MySQL

Extensions & New Features

- Multi-source replication new in 10.0
- GTID new in 10.0, Global Transaction ID
- SHOW EXPLAIN new in 10.0, gives the EXPLAIN query plan
- Roles new in 10.0
- PCRE new in 10.0, **P**erl **C**ompatible **R**egular **E**xpressions
- DELETE ... RETURNING new in 10.0

<https://mariadb.com/kb/en/library/mariadb-vs-mysql-features/>

MariaDB vs. MySQL

Truly Open Source

- All code in MariaDB : released under GPL, LGPL or BSD
- MariaDB client libraries (for C, for Java (JDBC), for Windows (ODBC) are released under **LGPL**
 - allows link with closed source software
 - MySQL client libraries are released under GPL that does not allow linking with closed source software
- MariaDB includes test cases for all fixed bugs.
 - Oracle doesn't provide test cases for new bugs fixed in MySQL 5.5
 - All bugs / development plans are public
 - <https://jira.mariadb.org/>
 - <https://github.com/MariaDB>

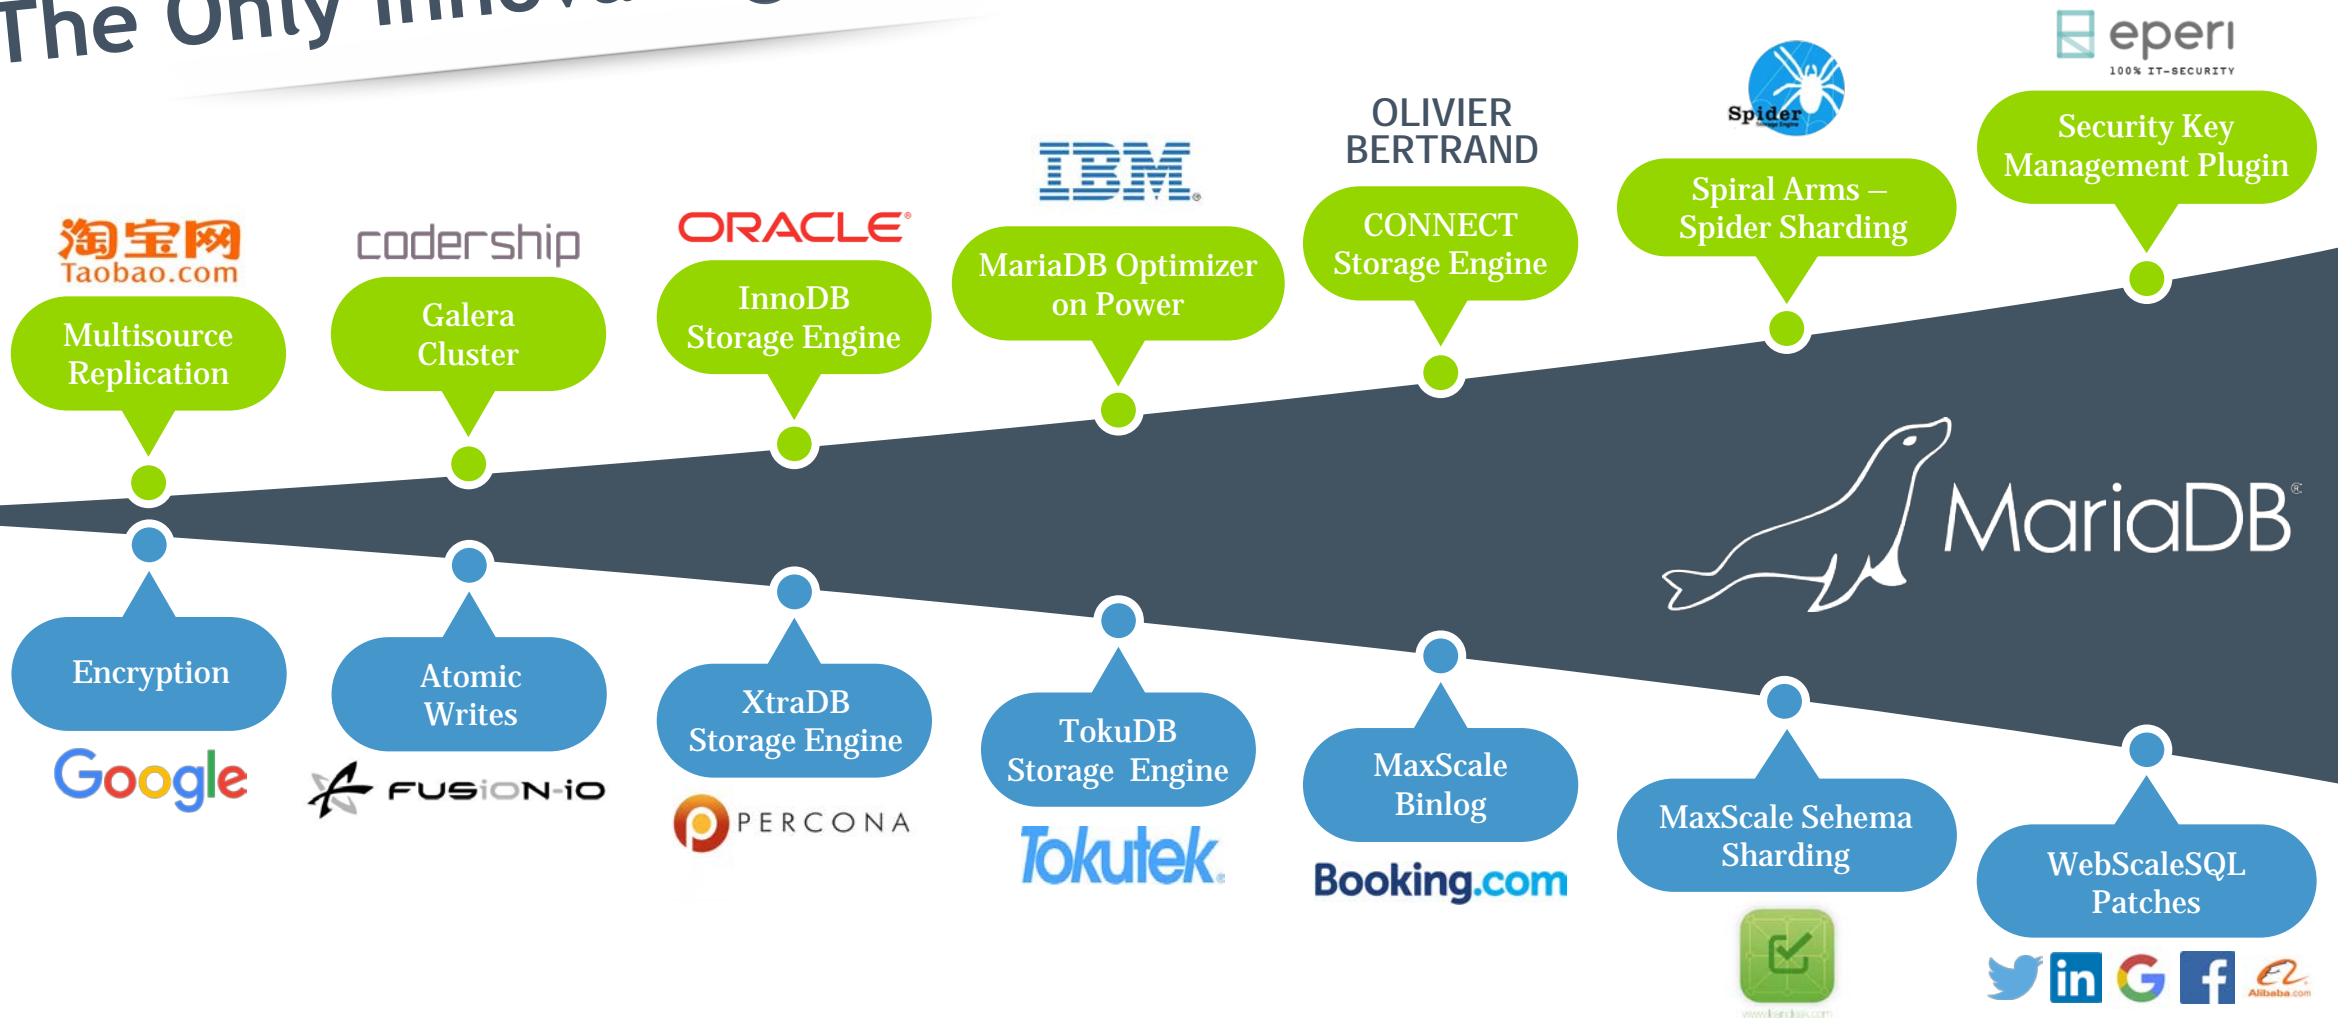


“

By 2018, more than 70% of new in-house applications will be developed on an ODBMS, and 50% of existing commercial RDBMS instances will have been converted or will be in process.

Gartner | STRATEGIC PLANNING ASSUMPTION

The Only Innovating Community



A decorative network graph pattern is located in the bottom-left corner of the slide. It consists of numerous small white and blue dots connected by thin white lines, forming a complex, web-like structure that extends towards the center of the slide.

We're building a database that is easy to use, easy to extend, and easy to deploy:

on premise, in the cloud, or hybrid, operational or analytical
— and with the languages and frameworks you prefer.


Why consider MariaDB?

Reduced Costs


Annual Subscriptions

Cloud Infrastructure

Modern Hardware



Economics of Inevitable Change



On premise...

- Oracle costs **80x** more
- Organizations can save **\$9 million**

In the cloud...

- On AWS, Oracle costs **145x** more
- On Oracle Cloud, it costs **69x** more

Default Database on Leading Linux Distro, Available on Leading Cloud Platforms

Linux Distributions



Cloud Services & Stacks



MariaDB Going Mainstream

OpenStack user survey
April 2017

Which databases are used for OpenStack components?

MariaDB Galera Cluster moved from fourth to first place on this list, up 7 points, while MongoDB slipped 6 points. MySQL with Galera also dropped significantly, down 8 points. This follows a longer-

term trend over time where deployments are moving away from MySQL (down 22% over past 18 months) to MariaDB and others. IBM DB2 was noted among other databases used.





Figure 5.7 n=478



MariaDB Going Mainstream



























Distribution	Estimated Reach	MariaDB Default since
CentOS	10 million	2013
RHEL	5.3 million	2013
openSUSE	28 million	2015
Fedora	13.4 million	2015
Debian	1.5 million	2017
Ubuntu	25 million	Expected to follow Debian
BSD, Mageia, macOS etc	1 million	Mostly all default MariaDB

MariaDB Going Mainstream

DB-Engines Ranking RDBMS Sep 2017

db-engines.com/en/ranking

137 systems in ranking, September 2017

Rank			DBMS	Database Model	Score		
Sep 2017	Aug 2017	Sep 2016			Sep 2017	Aug 2017	Sep 2016
1.	1.	1.	Oracle  	Relational DBMS	1359.09	-8.78	-66.47
2.	2.	2.	MySQL  	Relational DBMS	1312.61	-27.69	-41.41
3.	3.	3.	Microsoft SQL Server  	Relational DBMS	1212.54	-12.93	+0.99
4.	4.	4.	PostgreSQL  	Relational DBMS	372.36	+2.60	+56.01
5.	5.	5.	DB2 	Relational DBMS	198.34	+0.87	+17.15
6.	6.	6.	Microsoft Access	Relational DBMS	128.81	+1.78	+5.50
7.	7.	7.	SQLite	Relational DBMS	112.04	+1.19	+3.41
8.	8.	8.	Teradata	Relational DBMS	80.91	+1.67	+7.84
9.	9.	9.	SAP Adaptive Server	Relational DBMS	66.75	-0.16	-2.41
10.	10.	10.	FileMaker	Relational DBMS	61.00	+1.35	+5.64
 11.	11.	 13.	MariaDB 	Relational DBMS	55.47	+0.78	+16.94
12.	 13.	 11.	Hive 	Relational DBMS	48.62	+1.31	-0.21
13.	 12.	 12.	SAP HANA 	Relational DBMS	48.33	+0.36	+4.91
14.	14.	14.	Informix	Relational DBMS	27.84	+0.41	-0.35
15.	 16.	15.	Vertica 	Relational DBMS	22.01	+0.20	+0.95
16.	 15.	 17.	Microsoft Azure SQL Database	Relational DBMS	21.60	-0.31	+2.18
17.	17.	 16.	Netezza	Relational DBMS	19.40	-0.17	-0.41
18.	18.	18.	Firebird	Relational DBMS	17.86	-0.20	+2.28
19.	19.	 22.	Impala	Relational DBMS	13.56	+0.51	+4.34
20.	20.	 19.	Amazon Redshift 	Relational DBMS	13.04	+0.21	+2.21

MariaDB Going Mainstream

DB-Engines Ranking of RDBMS Sep 2017

db-engines.com/en/ranking





Financial Services



Travel



Retail & Ecommerce



Gvmt & Education



Telecom



Technology & Internet



Media & Social



12 Million Users in
45 Countries Trust Critical
Business Data to MariaDB

The database market is changing. It is time for a new leader.



What will this new leader provide?


Enterprise Reliability

High Availability (Replication / Failover)
Disaster Recovery (Backup & Restore)
Performance (Thread Pool)
Scalability (Partitioning and Sharding)
Security (Firewall and Encryption)

Open Source Innovation


Open Development (Tests and Roadmaps)
Community **Collaboration** (New Features)
Extensible Architecture (Custom Plugins)
Flexible Modeling (Dynamic Columns & JSON)
Streaming Integration (CDC and Kafka)

MariaDB Technology




SERVER

Enterprise-grade secure, highly available and scalable relational database with a modern, extensible architecture



MAXSCALE

Next-generation database **proxy** that manages security, scalability and high availability in scale-out deployments



COLUMNSTORE

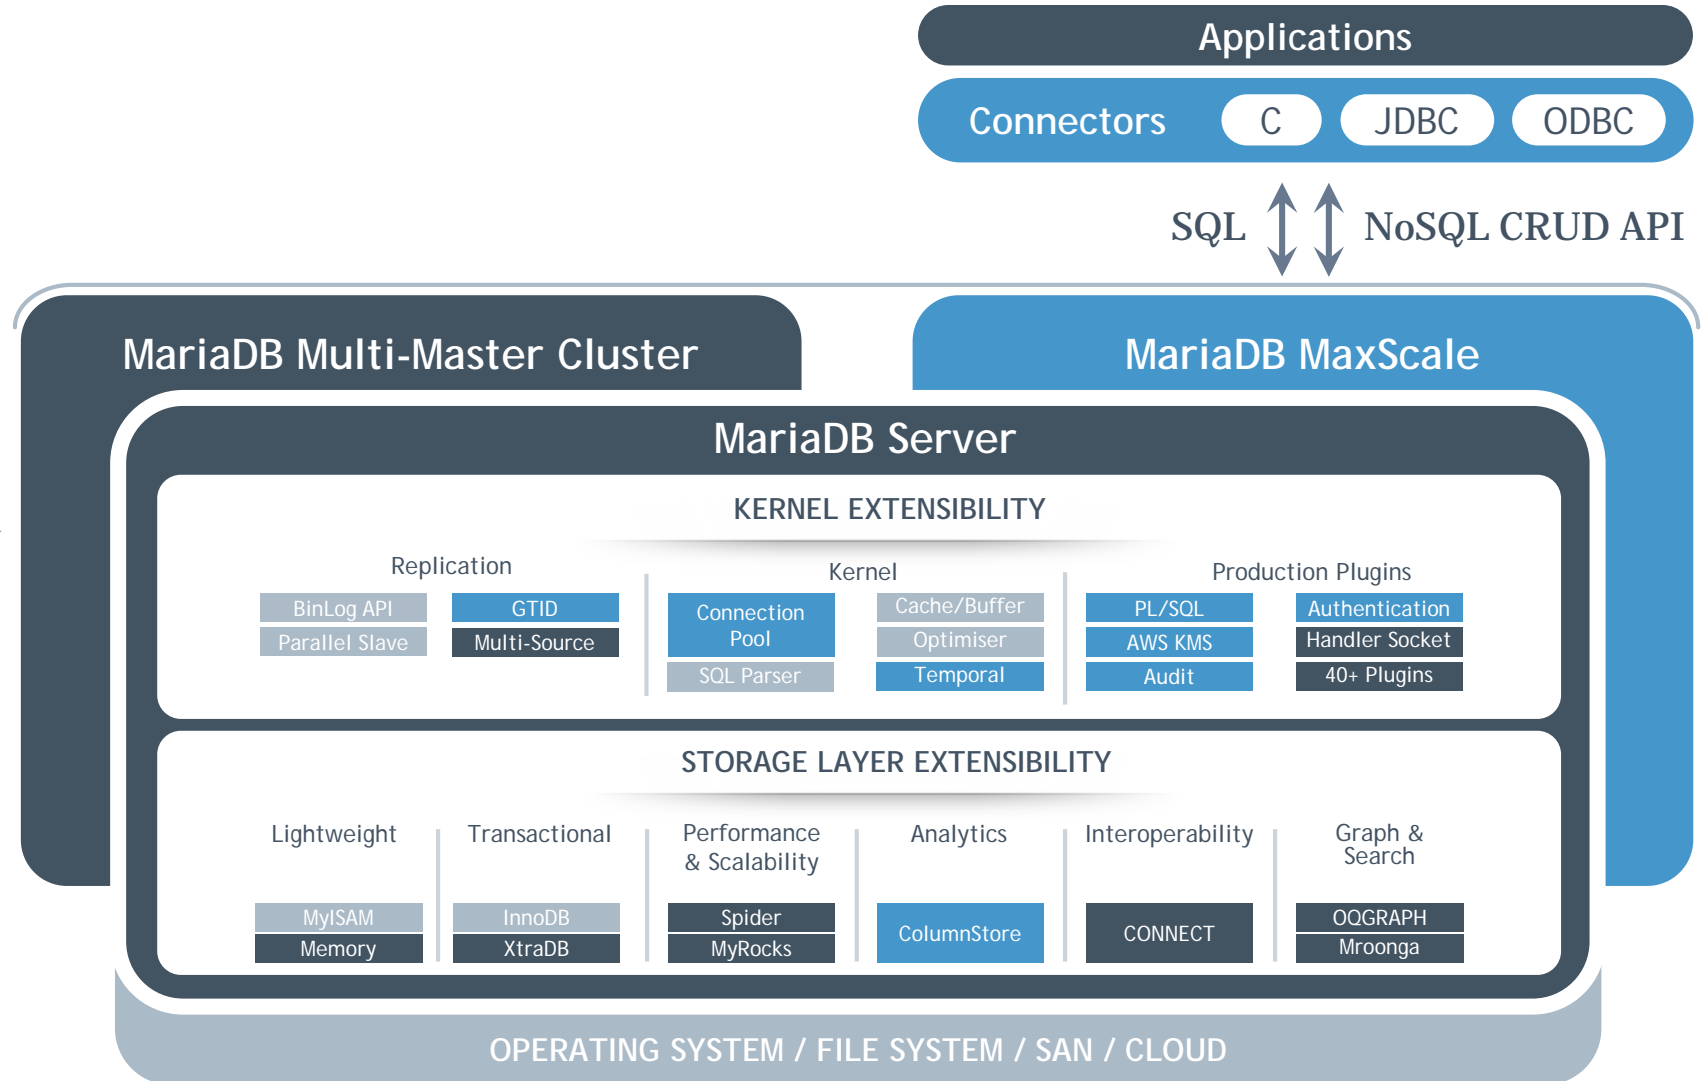
Columnar storage engine for massively parallel distributed query execution and data loading

Extensible Architecture

- Original Core
- MariaDB Engineering
- Community Contribution



MariaDB
Replicas
Supporting
Asynchronous,
Semi-Sync &
Synchronous
replication



MariaDB Storage Engines

MariaDB Storage Engines

General Purpose

- **XtraDB** : best choice in many cases until MariaDB 10.1
 –performance-enhanced fork of InnoDB and is MariaDB's default engine until MariaDB 10.1
- **InnoDB** : general transactional engine, default MariaDB **10.2**
- **MyISAM** : small footprint and allows for easy copying between systems
 –MySQL's oldest storage engine
- **Aria** : MariaDB's modern improvement on MyISAM

<https://mariadb.com/kb/en/library/choosing-the-right-storage-engine/>

MariaDB Storage Engines

Scaling, Partitioning

- Galera Cluster : synchronous multi-master cluster. **only InnoDB supported**
- **TokuDB** is a transactional storage engine which is optimized for workloads that do not fit in memory, and provides a **good compression ratio**
- **Spider** uses partitioning to provide data **sharding** through multiple servers
 - <https://www.slideshare.net/Kentoku/mariadb-103spider>
- **ColumnStore** utilizes a massively parallel **distributed** data architecture and is designed for big data scaling to process petabytes of data
 - <https://www.slideshare.net/InsightTechnology/dbtstky2017-c37-mariadb>

MariaDB Storage Engines

Compression / Archive

- **MyRocks** : enables greater compression than InnoDB, as well as less write amplification giving better endurance of flash storage(SSD) and improving overall throughput. Developed at **Facebook**
 - <https://enterprisezine.jp/dbonline/detail/9450>
- TokuDB : transactional storage engine which is optimized for workloads that do not fit in memory, and provides a good compression ratio
- Archive : best used for archiving

Storage Engines

Connecting to other data sources

- **CONNECT** : allows access to different kinds of text files and remote resources as if they were regular MariaDB tables
- CSV : read and append to files stored in CSV format.
 - since MariaDB 10.0, CONNECT is a better choice and is more flexibly able to read and write such files
- FederatedX : uses libmysql to talk to the data source, the data source being a remote RDBMS.
 - Currently, since FederatedX only uses libmysql, it can only talk to another MySQL RDBMS.
- CassandraSE : allowing access to an older version of Apache Cassandra NoSQL DBMS.
 - It was relatively experimental, and is **no longer being actively developed**

Storage Engines

Full Text Search

- **SphinxSE** : proxy to run statements on a remote Sphinx DB server
- **Mroonga** : fast CJK-ready full text search
 - <http://mroonga.org/ja/blog/2016/07/21/mariadb-community-event-in-tokyo.html>
 - <https://www.slideshare.net/kou/dbtechshowcasetokyo2017>

Storage Engines

Cache, read-only

- MEMORY : does not write data on-disk (all rows are lost on crash) and is best-used for **read-only caches** of data from other tables, or for temporary work areas.
 - With the default XtraDB and other storage engines having good caching, **there is less need for this engine than in the past.**

Evolution of MariaDB Server

10.1 (GA)

Multi-master Replication (Galera)

Data-at-Rest Encryption

Password Validation

Page Compression

Facebook Defragmentation

Spatial References (GIS)

10.2 (GA)

MyRocks

Flashback (Point in Time Rollback)

Common Table Expressions

Window Functions

Check Constraints

JSON & GeoJSON Functions

Multi-Trigger Support

Delayed Replication

Binary Log Compression

Per User Resource Limits

Virtual Column Indexes

10.3 (DEV)

Sequences & PL/SQL

System Versioned Tables

As Of (Point in Time Querying)

User-defined Aggregate Functions

Intersect & Except

Hidden Columns


MariaDB MaxScale

MariaDB MaxScale is a next-generation database proxy that manages security, scalability and high availability in a scale out deployment.



Security

Secure **database firewall** to prevent cyber attacks like SQL injection and DDoS




Scalability

Manage your scaled-out infrastructure **without changing application code**



High Availability

Ensure uptime with no single point of failure and **minimize downtime** during upgrade



Data Streaming

Stream transactional data to data lake for **real-time analytics**

MariaDB ColumnStore(AX)

High performance columnar storage engine that supports a wide variety of analytical use cases in highly scalable distributed environments

Better Price
Performance



Power of SQL and
Freedom of Open
Source to Big Data
Analytics

Easier Enterprise
Analytics



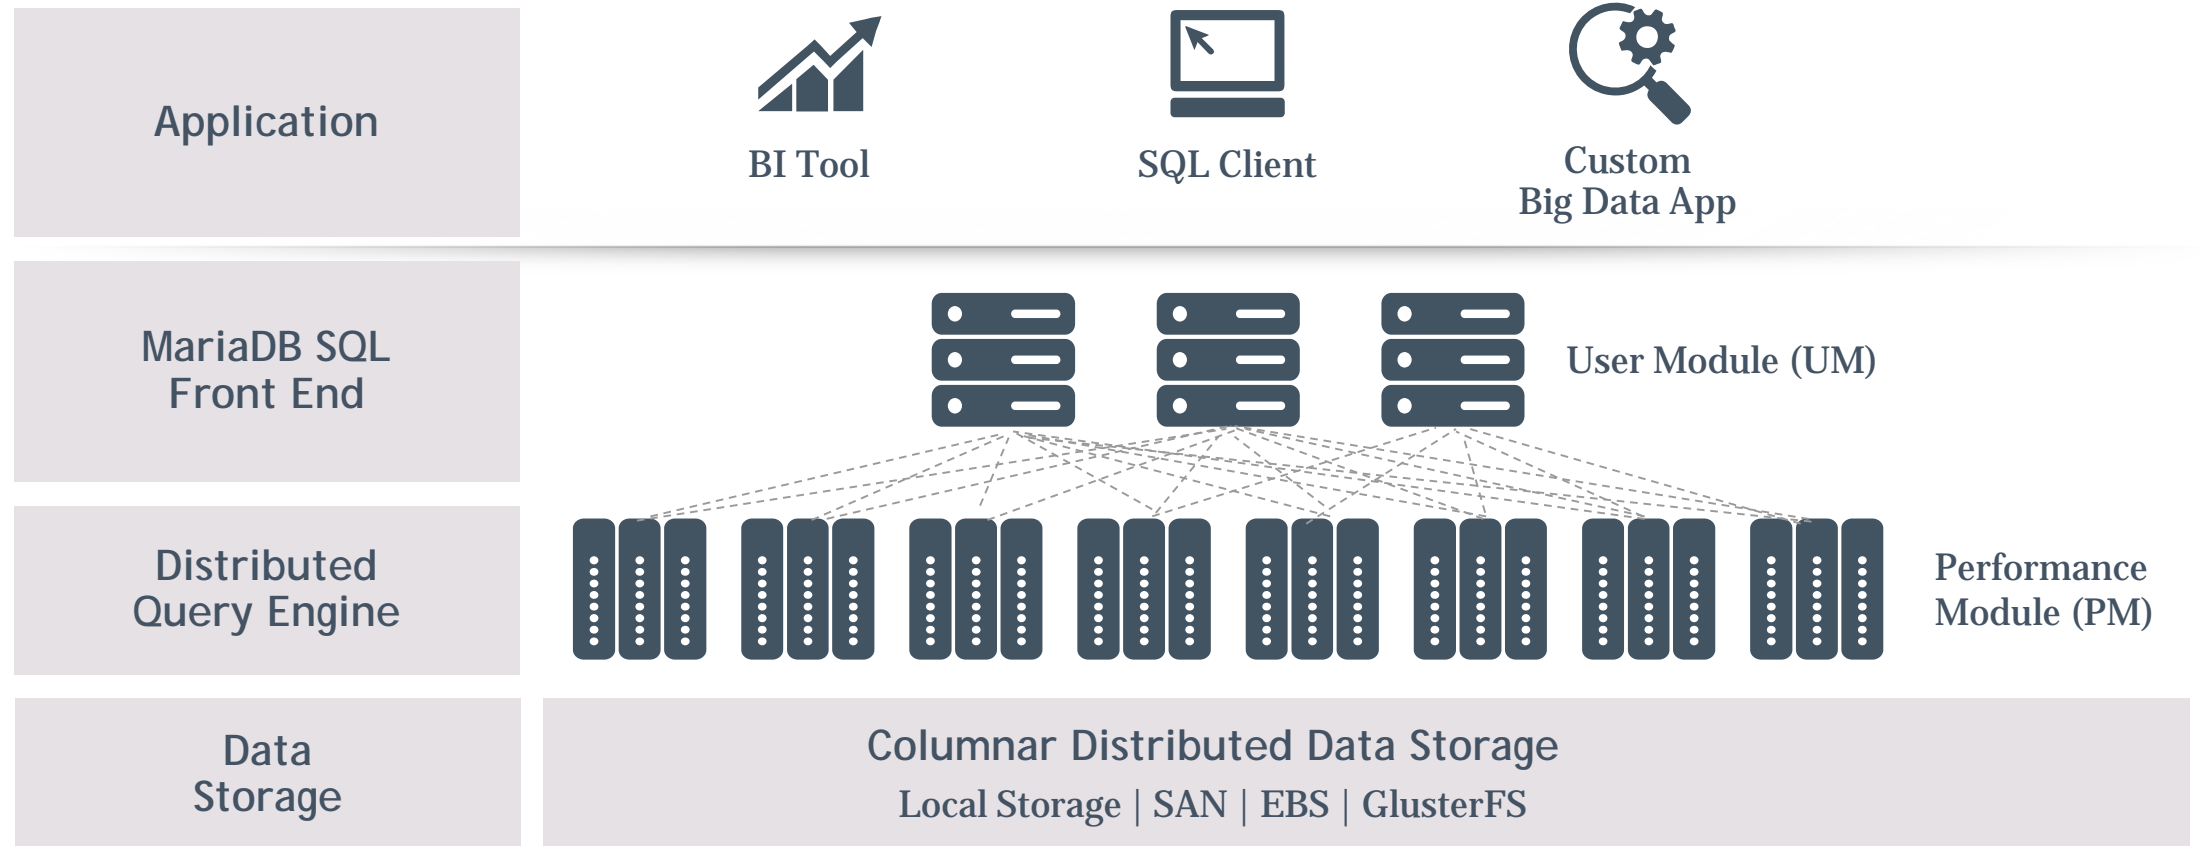
Single Interface for
OLTP and analytics
Easy to Manage and
Scale

Faster, More
Efficient Queries



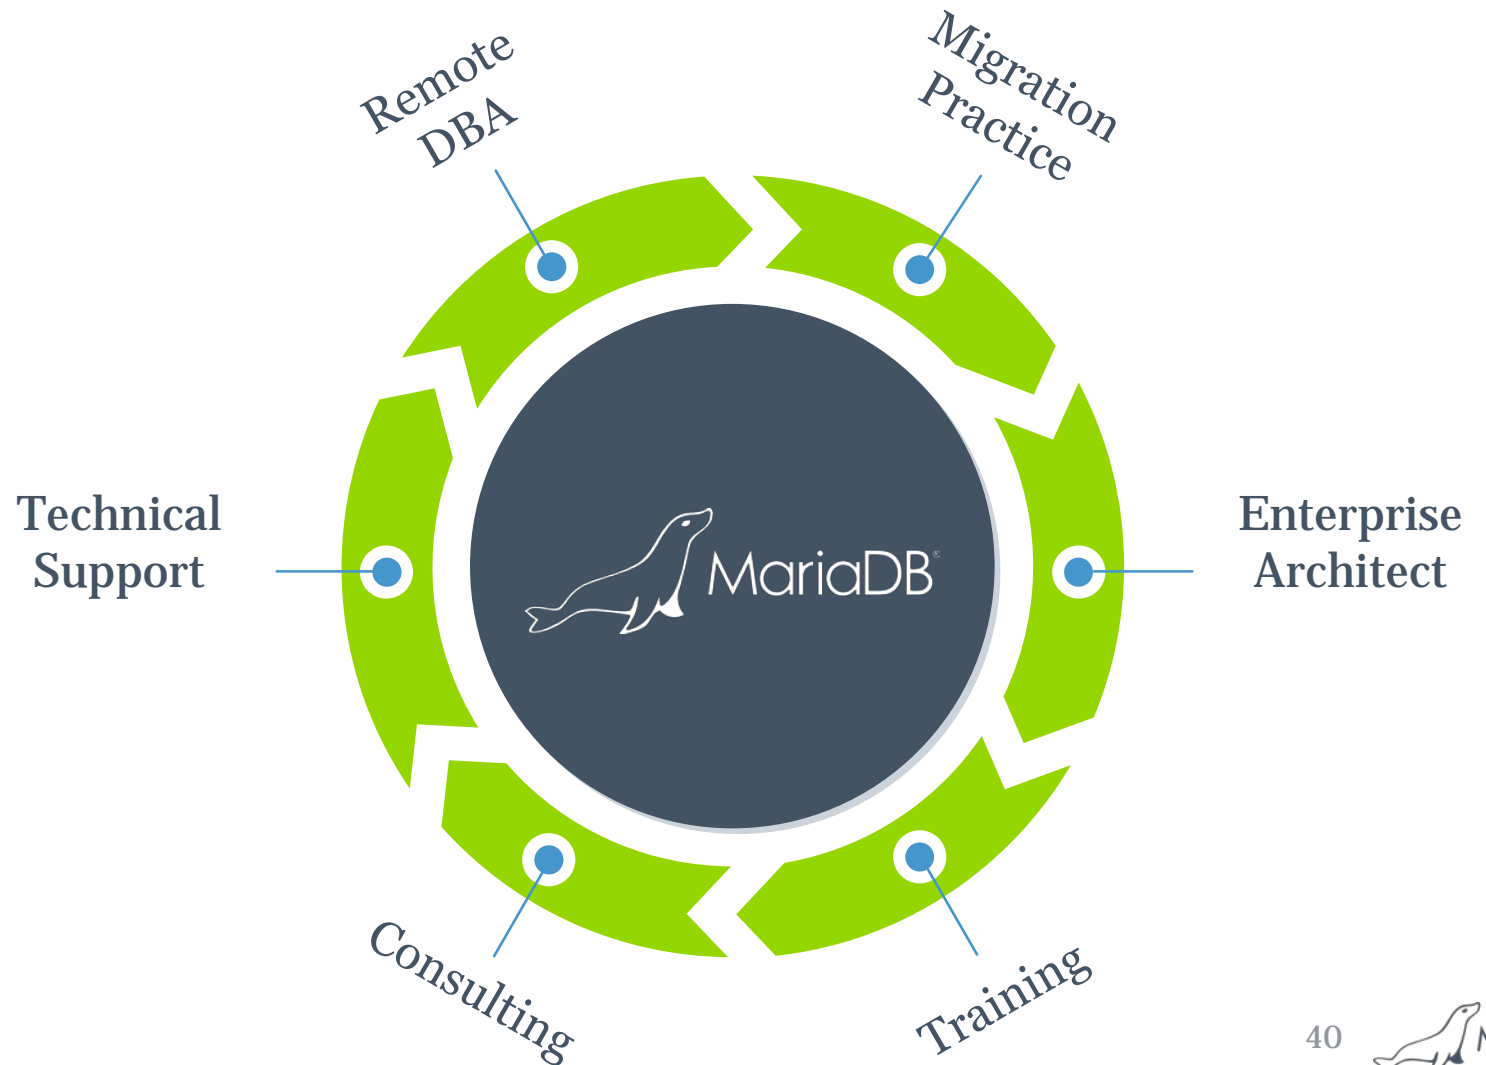
Parallel query
processing for distributed
environments

MariaDB ColumnStore Architecture



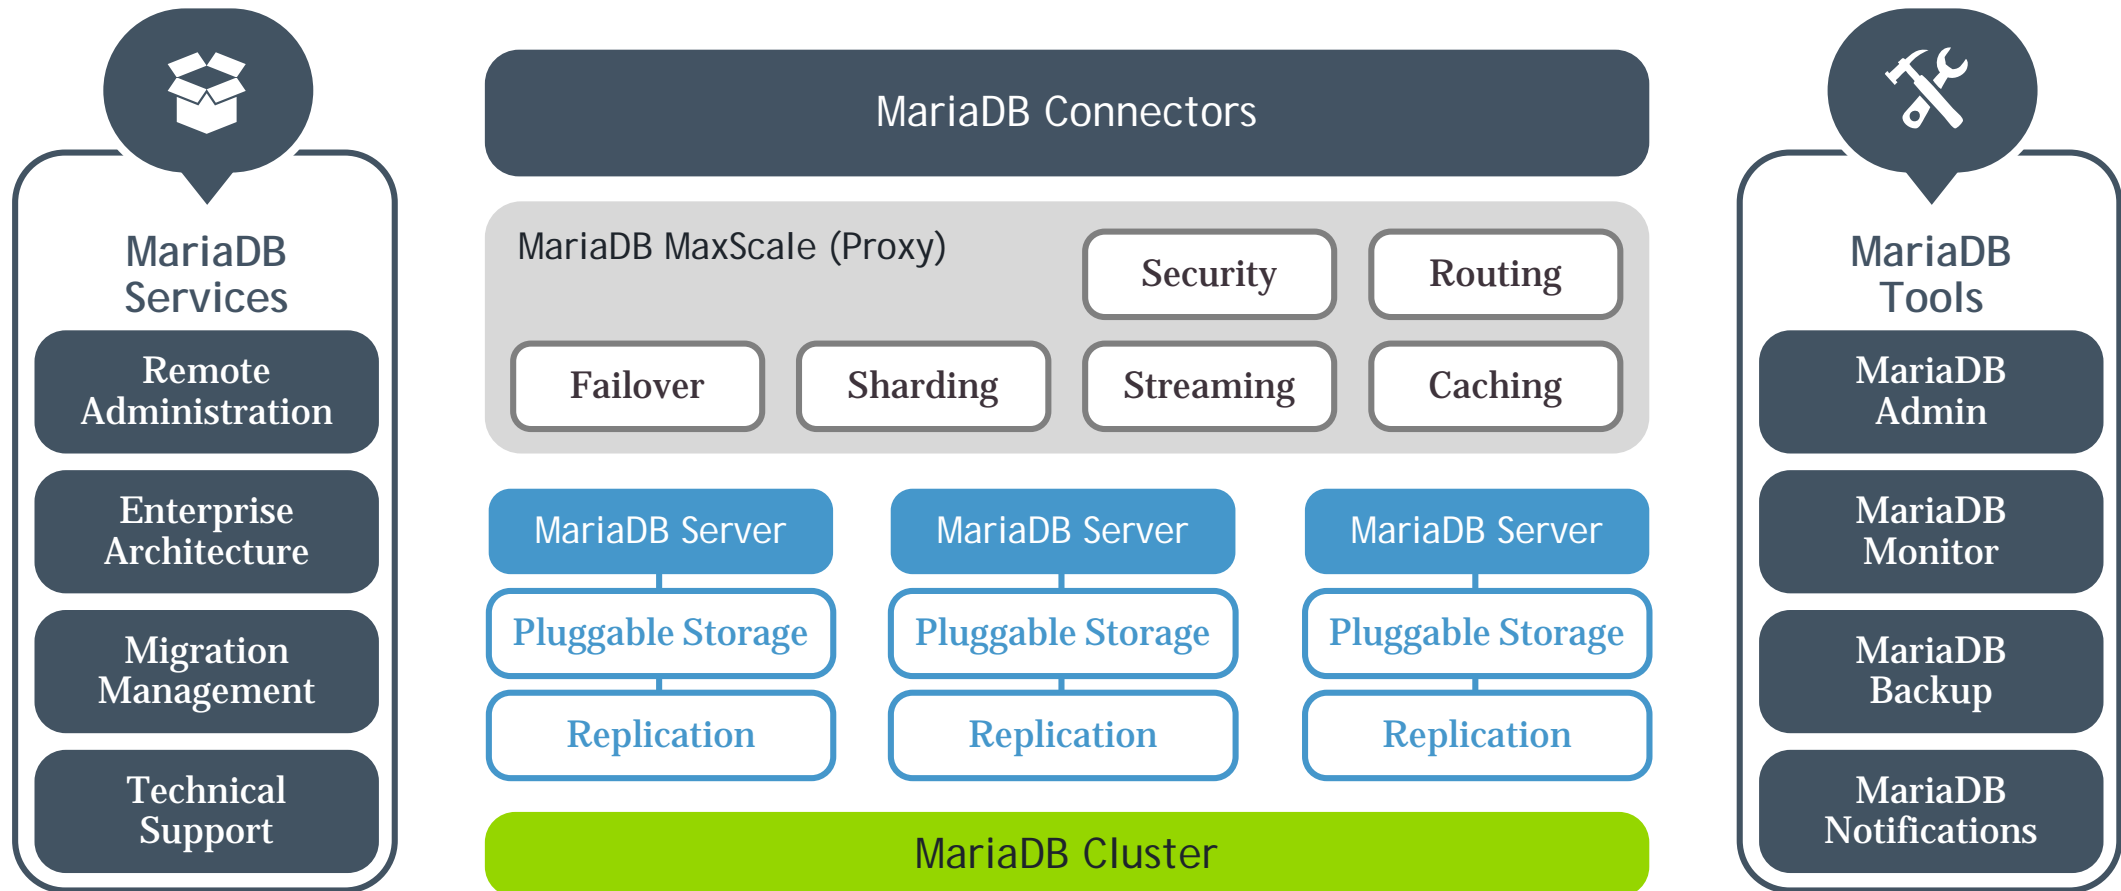
Services and Support

Put our
Expertise
to Work
for You



MariaDB TX 2.0

Transaction Platform



MariaDB TX 2.0

Transaction Platform



Software

licensing
and support for
databases, the
database proxies and
database connectors



Services

credits for remote
administration,
enterprise architecture,
migration planning
services and more



Tools

tools for
administration,
monitoring, backup
and replication
management

Customer and Use Cases



- Multi-terabyte DB
- 80M transactions / month



- 250 servers, 600G + 1.5T archive
- 10M travelers/quarter
- 4M transactions/ month



WIKIPEDIA
The Free Encyclopedia

- ~14TB in MariaDB production clusters



- 50+ Node Cluster
- Multi-billion rows
- 600 Million reads/second



- Over 150 servers
- 150-200k queries / sec on the MariaDB Cluster



- 3 to 10 TB
- Over billion rows, most tables 100's of millions of rows



- 70 million rows per day
- 4 billion impressions per month



- Over 5 TB in Pay Per click application



- 6TB and millions of CDR's

Get Started with MariaDB

1

Download

<https://mariadb.com/downloads>

2

Read the Technical overviews

<https://mariadb.com/resources/datasheets-guides>

3

Search the Knowledge Base

<https://mariadb.com/kb>

4

Watch a Webinar

<https://mariadb.com/resources/webinars>

Upcoming events





Home / Conferences/Events / Community /
Development / 2017-2 Developers Unconference and
Related Events, Shenzhen

2017-2 Developers Unconference and Related Events, Shenzhen



2017-08-02

Written by Ian

4 Comments

Gilfillan

Platinum Sponsors

 Alibaba Cloud Tencent Cloud

13-17
November,
2017

M|18

FEB 26-27, 2018

CONRAD NEW YORK

m18.mariadb.com



Latest news



What is new in ColumnStore 1.1 ?

- Data Engine
 - Columnar Engine based on MariaDB 10.2
- Streaming
 - Native Data API for ColumnStore files: C++: LGPL
 - Data Adapters
 - Avro, JSON - BSL
 - Streaming Change Data Capture via MaxScale & Kafka - BSL
 - Streaming Insert via MaxScale - BSL
- HA
 - Built-in Data Redundancy/ Integrated **GlusterFS** for Data HA
- Analytics
 - User defined distributed aggregate and Window functions
- Data Types
 - Text, BLOB
- Ease of Use
 - Backup, restore tool - BSL
- Performance
 - Improved string handling / memory utilization
 - General performance improvements
- Security
 - Audit Plugin Integration
- Certification
 - Tableau



Thank you